



# An (extremely brief) introduction to Bayesian Markov chain Monte Carlo

Matt Denwood

**md@sund.ku.dk**

Department of Veterinary and Animal Sciences

University of Copenhagen

Denmark



# Overview

- What is Bayesian statistics?
- What is the theory behind MCMC?
- How does MCMC work in practice?
- A basic example: prevalence estimation



# What is Bayesian statistics?



# Bayes theorem

Discrete form:

$$P(A_j | B) = \frac{P(B | A_j) \cdot P(A_j)}{\sum_i P(B | A_i) \cdot P(A_i)}$$

Continuous form:

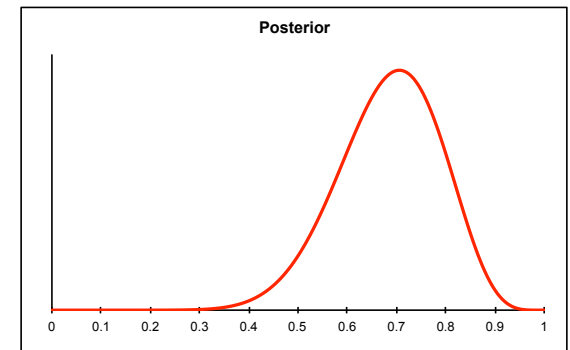
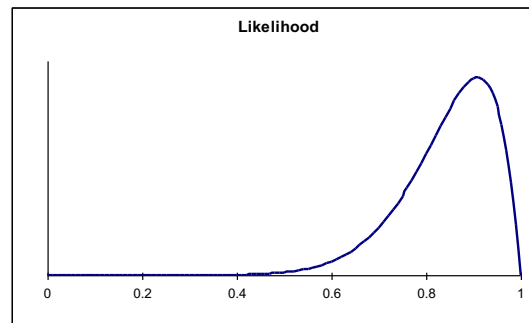
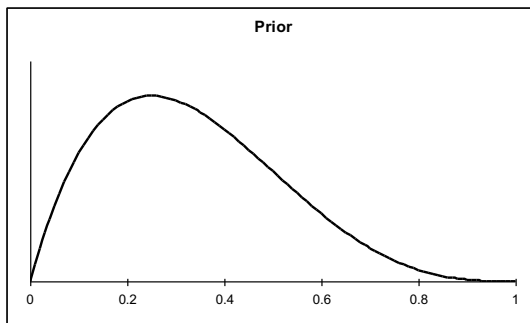
$$f_X(x | Y = y) = \frac{f_Y(y | X = x) \cdot f_X(x)}{\int_{-\infty}^{\infty} f_Y(y | X = \varepsilon) \cdot f_X(\varepsilon) \cdot d\varepsilon}$$

General form:

Posterior  $\propto$  Likelihood  $\times$  Prior

# Bayesian statistics

- "Belief" is expressed as a probability distribution
  - i.e. We don't know exactly what the parameter value is, but we can quantify our uncertainty
  - This is the key difference to frequentist statistics
- We usually have continuous probability distributions for both prior and posterior



What we knew before + What the data tell us = What we know now

# What is the theory behind MCMC?



# MCMC: the theory

First define a function to calculate a posterior probability corresponding to a Binomial likelihood and Beta prior:

```
log_posterior_fun <- function(parameter){  
  # ... Some R code e.g. ...  
  ll <- dbinom(data$Pos, data$N, parameter, log=TRUE)  
  lp <- dbeta(parameter, 1, 1, log=TRUE)  
  return(ll + lp)  
}
```

Then choose a 'reasonable' place to start looking for our parameter values, and calculate the posterior at that parameter value:

```
parameter[1] <- 0.25  
log_post[1] <- log_posterior_fun(parameter[1])  
log_post[1]  
-27.5222
```



# MCMC: the theory

Now sample another parameter value using a random walk:

```
new_par <- rnorm(1, mean=parameter[1], sd=sigma)
          0.2606076
new_lpost <- log_posterior_fun(new_par)
          -18.00832
```

That's a bit better! Let's store this new value:

```
parameter[2] <- new_par
log_post[2] <- new_lpost
```

And resample:

```
new_par <- rnorm(1, mean=parameter[2], sd=sigma)
          0.2572526
new_lpost <- log_posterior_fun(new_par)
          -20.72975
```





## MCMC: the theory

This is a bit worse. Should we keep it or not?

We will let fate decide ... but it makes sense to have a higher chance of accepting the new parameter value if it is not much worse than the old value:

```
probability_ratio <- exp(new_lpost - log_post[2])  
0.06578108  
accept <- rbinom(1, 1, probability_ratio)  
0
```

So we reject this new value, and keep the old one

We also reward the old parameter for being hard to beat by counting it twice (so this counts as the 2<sup>nd</sup> *and* 3<sup>rd</sup> values in the chain)

```
parameter[3] <- parameter[2]  
log_post[3] <- log_post[2]
```

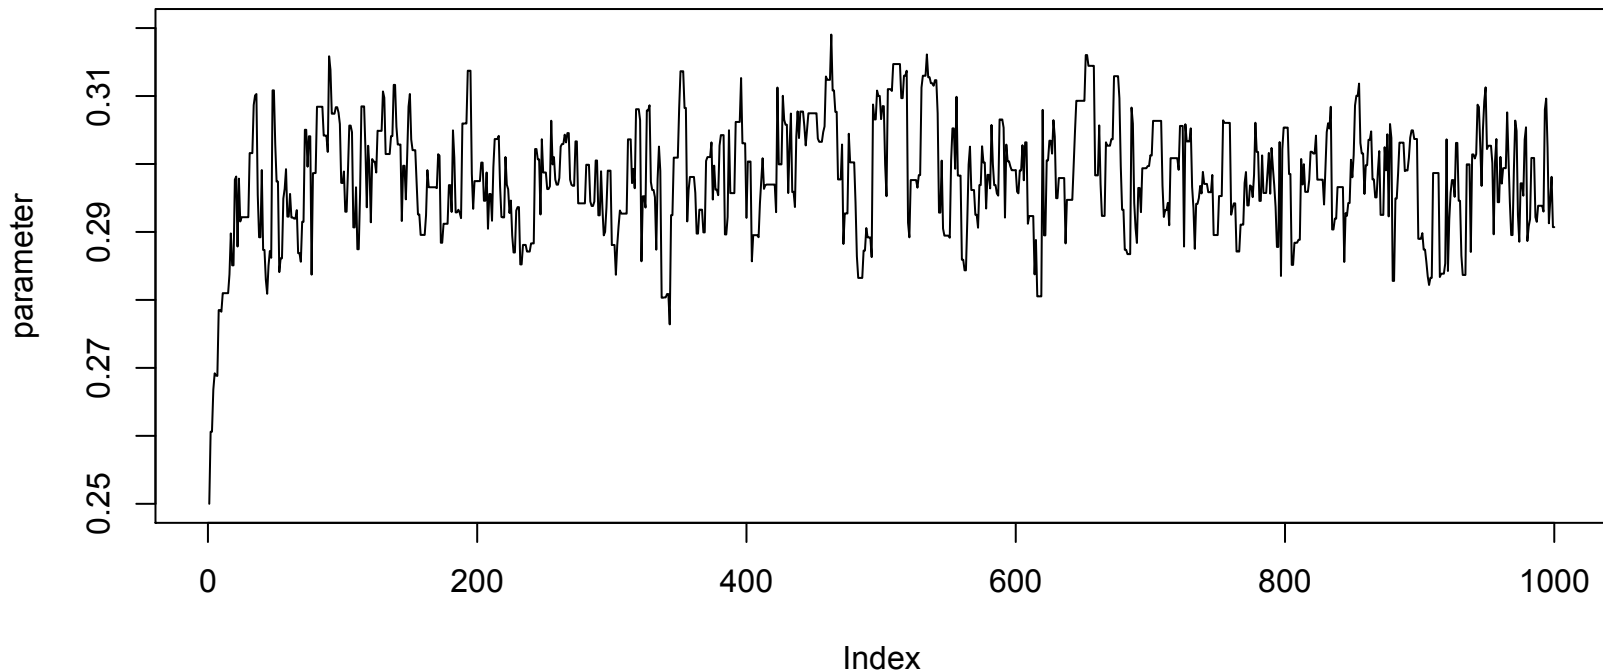


# MCMC: the theory

Repeat for the rest of the 1000 iterations, then look at the results:

```
plot(parameter, type='l')
```

Trace plot of parameter values



# Markov chain Monte Carlo

This algorithm is a basic implementation of MCMC!

- MCMC is used to approximate a posterior distribution by sampling from it
- Dependence on previous value is critical
  - Consecutive samples are *not* independent
  - Eventually a Markov chain will forget where it started, but it might take a while before it does this

This non-randomness of Markov chains has 2 important consequences:

1. The starting values are not immediately forgotten
  - We need a *burn-in* period before sampling from the chains (aka identifying convergence on the stationary distribution)
2. A sample of 1000 auto-correlated iterations gives us less information than a sample of 1000 independent iterations
  - We need to know the *effective sample size* or *Monte Carlo error*

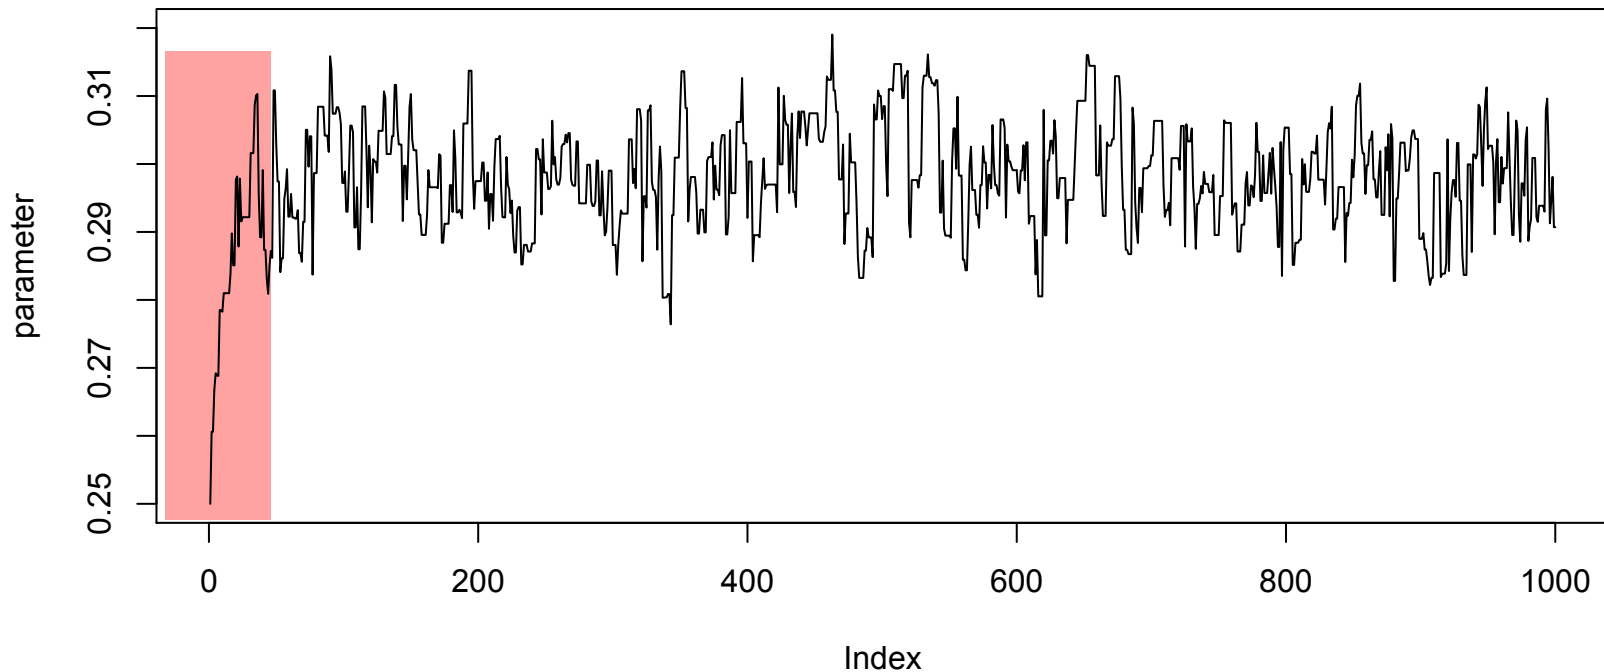


# 1. Burn-in

It is ESSENTIAL to identify and remove the burn-in period

```
plot(parameter, type='l')
```

Trace plot of parameter values

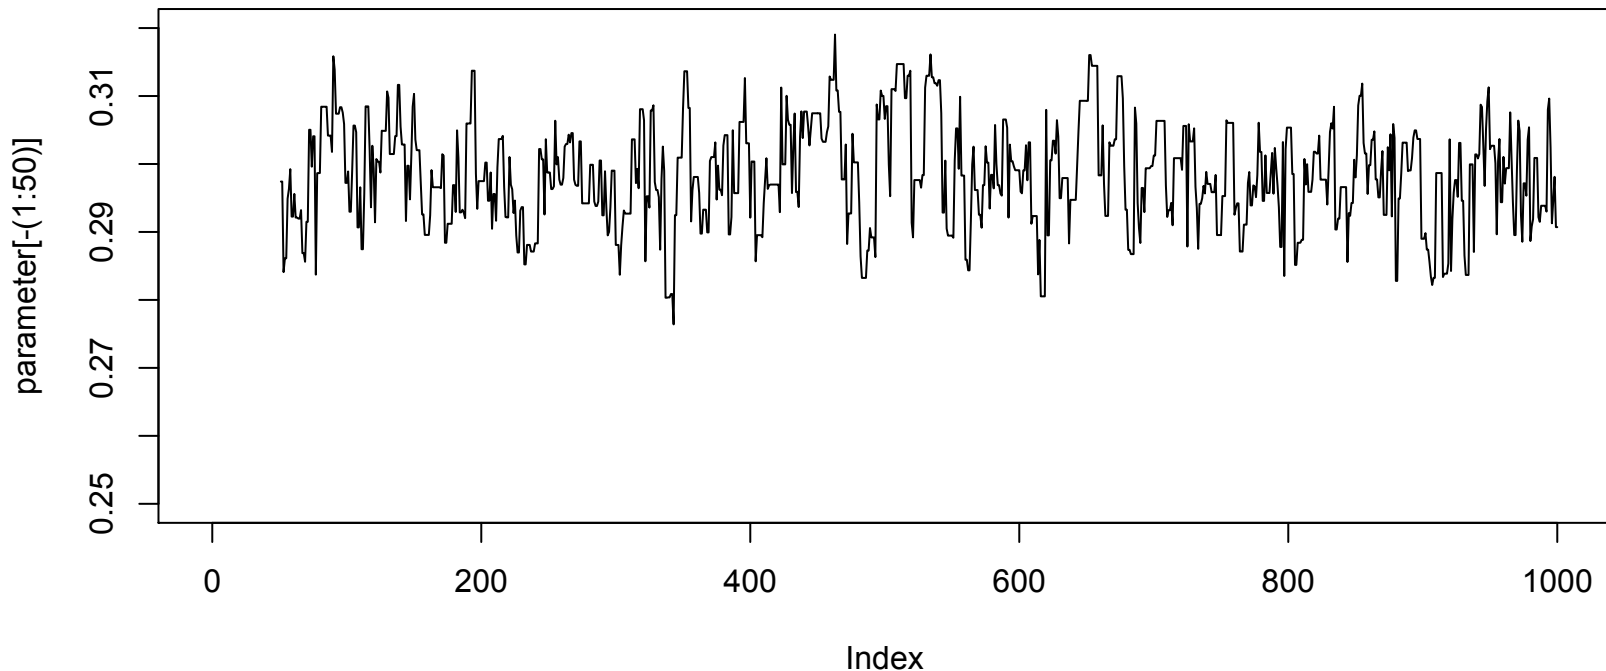


# 1. Burn-in

It is ESSENTIAL to identify and remove the burn-in period

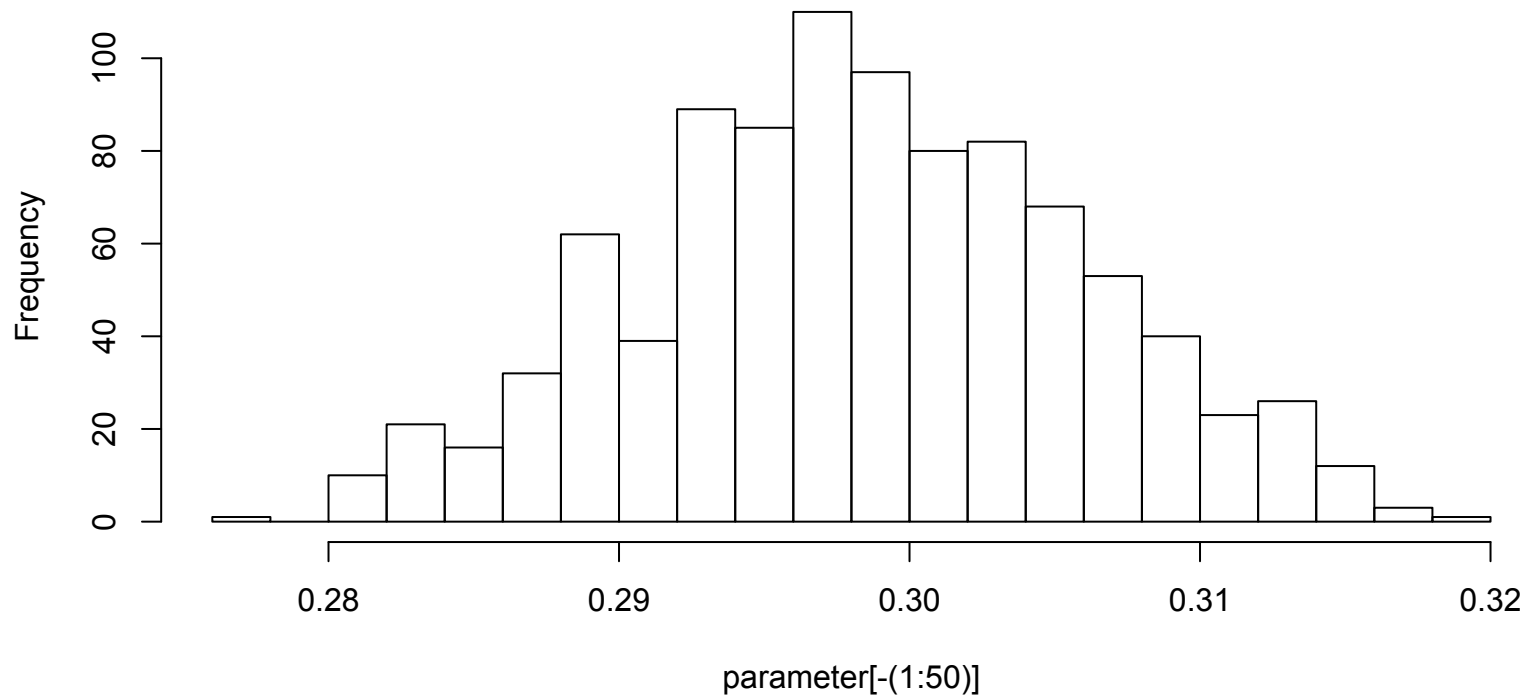
```
plot(parameter[-(1:50)], type='l')
```

Trace plot of stationary chain



## 2. Effective sample size

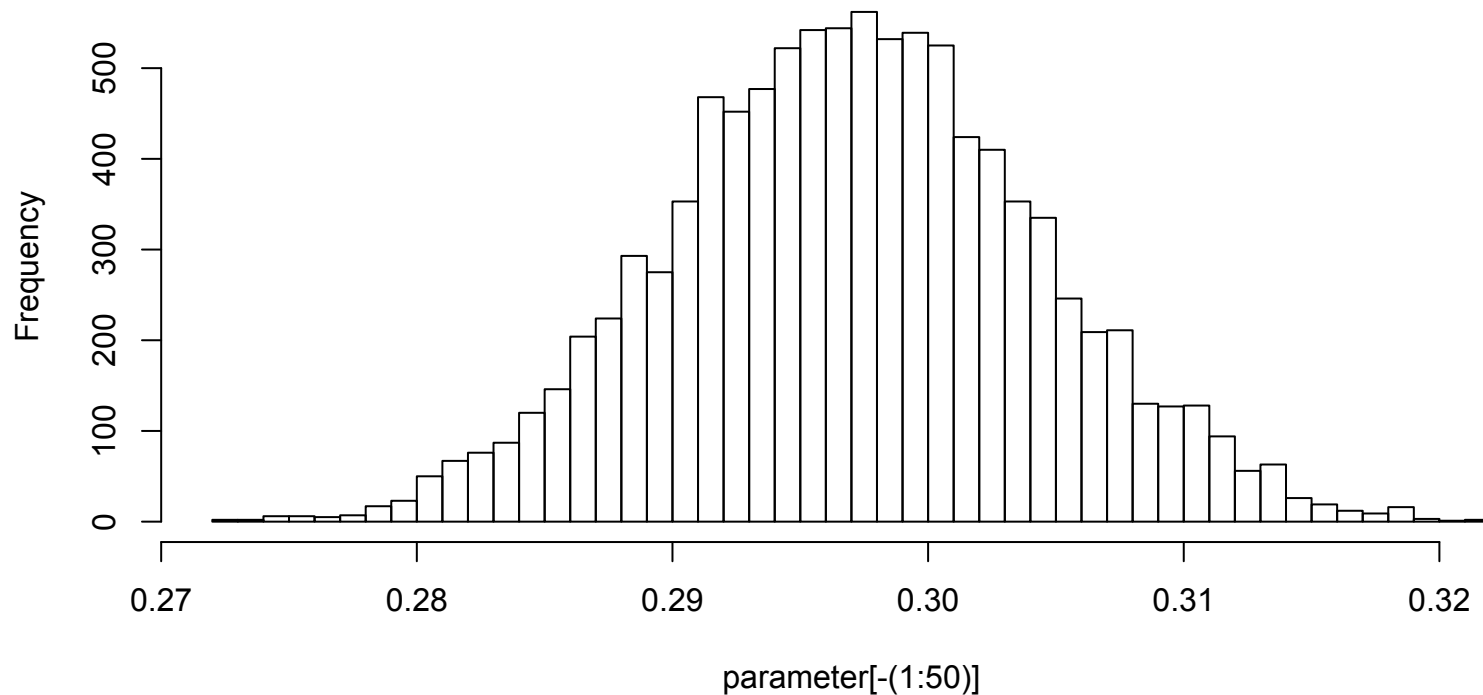
Histogram of 950 sampled values



An effective sample size of 137 – not a good approximation

## 2. Effective sample size

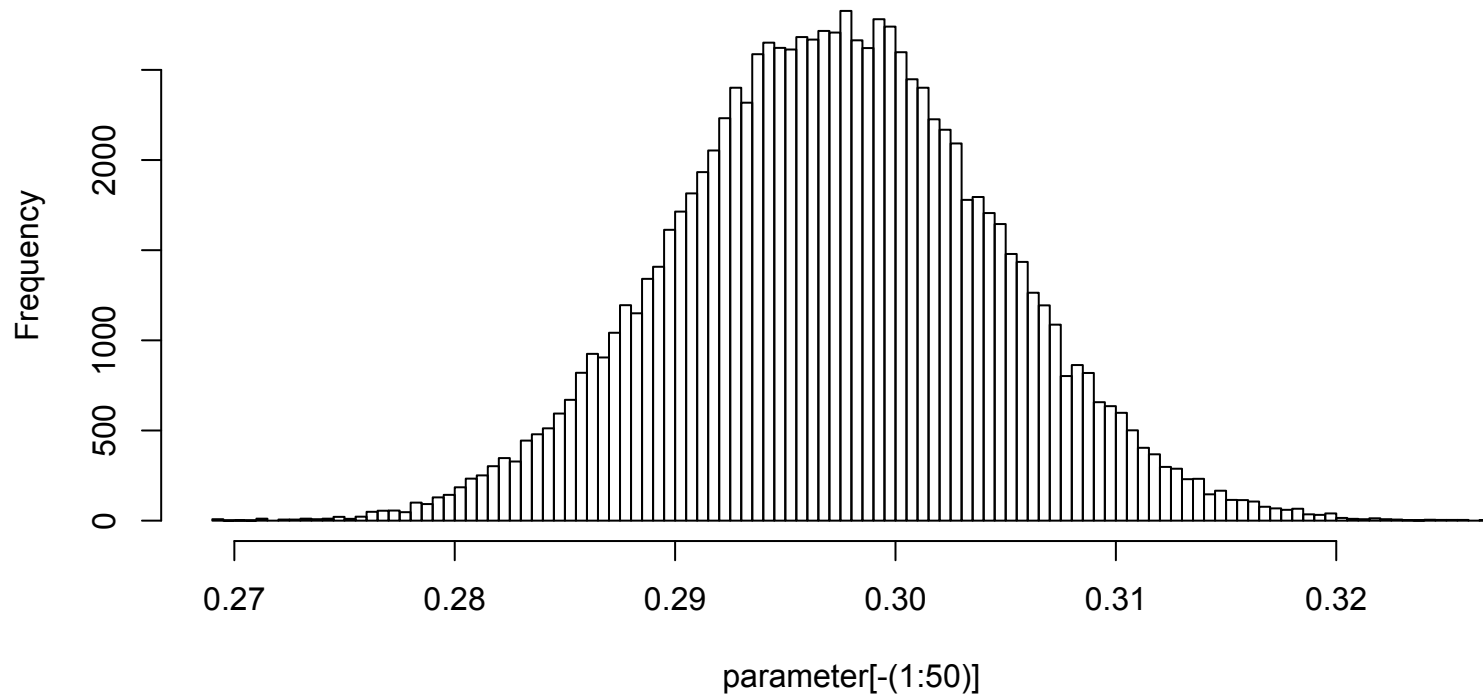
Histogram of 10000 sampled values



An effective sample size of 1752 – a pretty good approximation

## 2. Effective sample size

Histogram of 100k sampled values



An effective sample size of 16759 – a very good approximation



# Dangers of MCMC

1. Problems related to the Markov chain
  - We didn't remove enough burn-in iterations so our posterior is skewed by the parameter space from where we happened to start the initial values
  - Our chain hasn't really found the stationary posterior, and is just sampling from a local solution
  - **CONVERGENCE**
2. Problems related to the Monte Carlo integration
  - We didn't take enough samples to adequately represent the true posterior
  - **EFFECTIVE SAMPLE SIZE / MONTE CARLO ERROR**



# How does MCMC work in practice?



# Using BUGS

## Bayesian (analysis) Using Gibbs Sampling

- ✓ Generic MCMC framework with ready-made algorithms
- ❖ Does not automatically resolve any of the dangers with MCMC sampling!

A BUGS model is compiled by the software into an MCMC sampler

- All variables in the model must be defined EXACTLY once
- NB: it is a symbolic language NOT a programming language

## OpenBUGS

- Available for Windows, Linux and Mac (under emulation)
- No longer under active development

## JAGS (Just Another Gibbs Sampler)

- Technically not BUGS, but almost identical
- Primarily designed to be called from within R
  - Packages available: rjags, runjags, R2jags, jagsUI



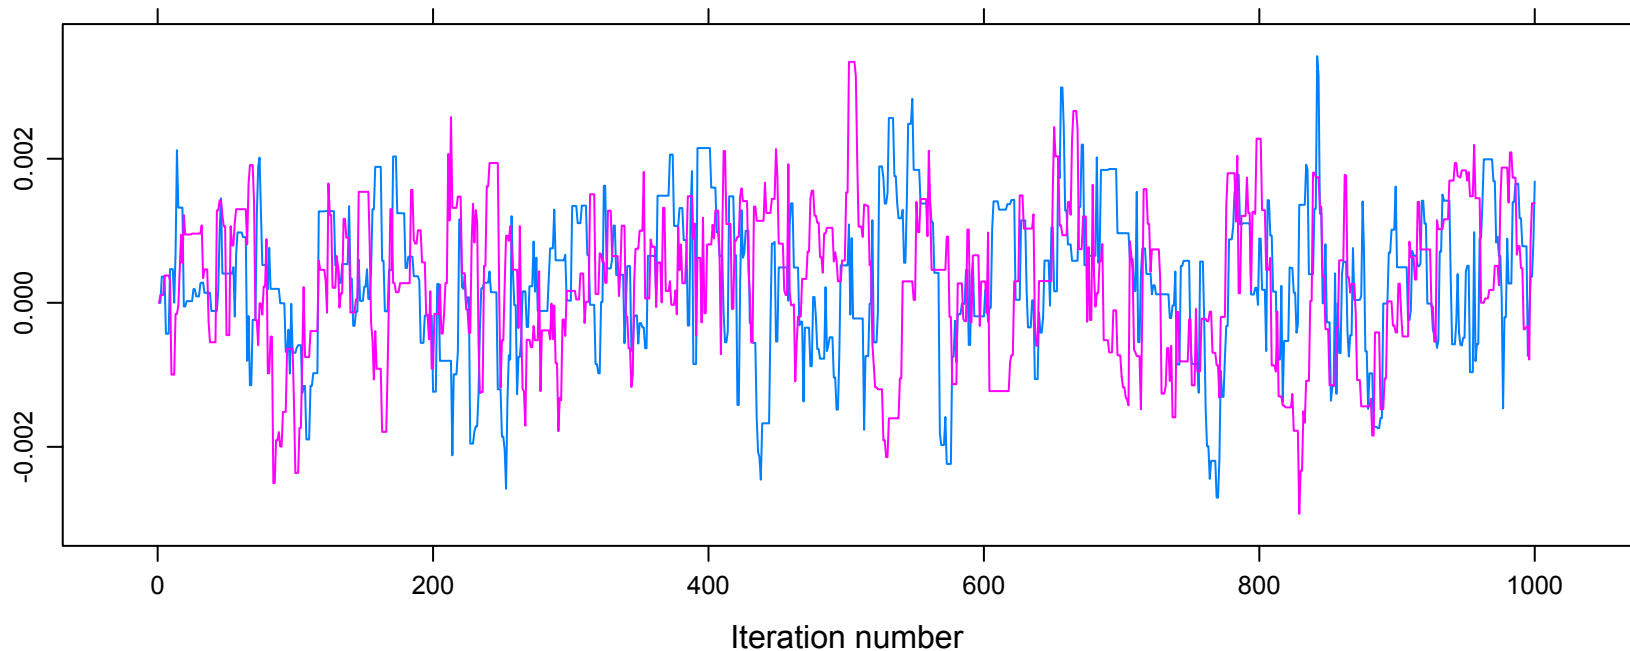
# Procedure

- 1) Define a syntactically valid model (used to calculate the posterior)
- 2) Supply data and initial values for at least 2 chains (these are often optional)
- 3) Decide which variables to monitor
- 4) Run the model
- 5) Make sure the chains have converged
  - Brooks-Gelman-Rubin statistic compares the variance within chains to the variance between chains – ratio  $< 1.05$  indicates that chains are sampling the same values so have probably converged
  - But always check the trace plots visually as well!
- 6) Make sure the number of sampled iterations is high enough
  - Minimum effective sample size of  $>400$  for parameters of interest
  - Or equivalently when  $MC\_error < 0.05 \times \text{sample SD}$



# Visual assessment of convergence

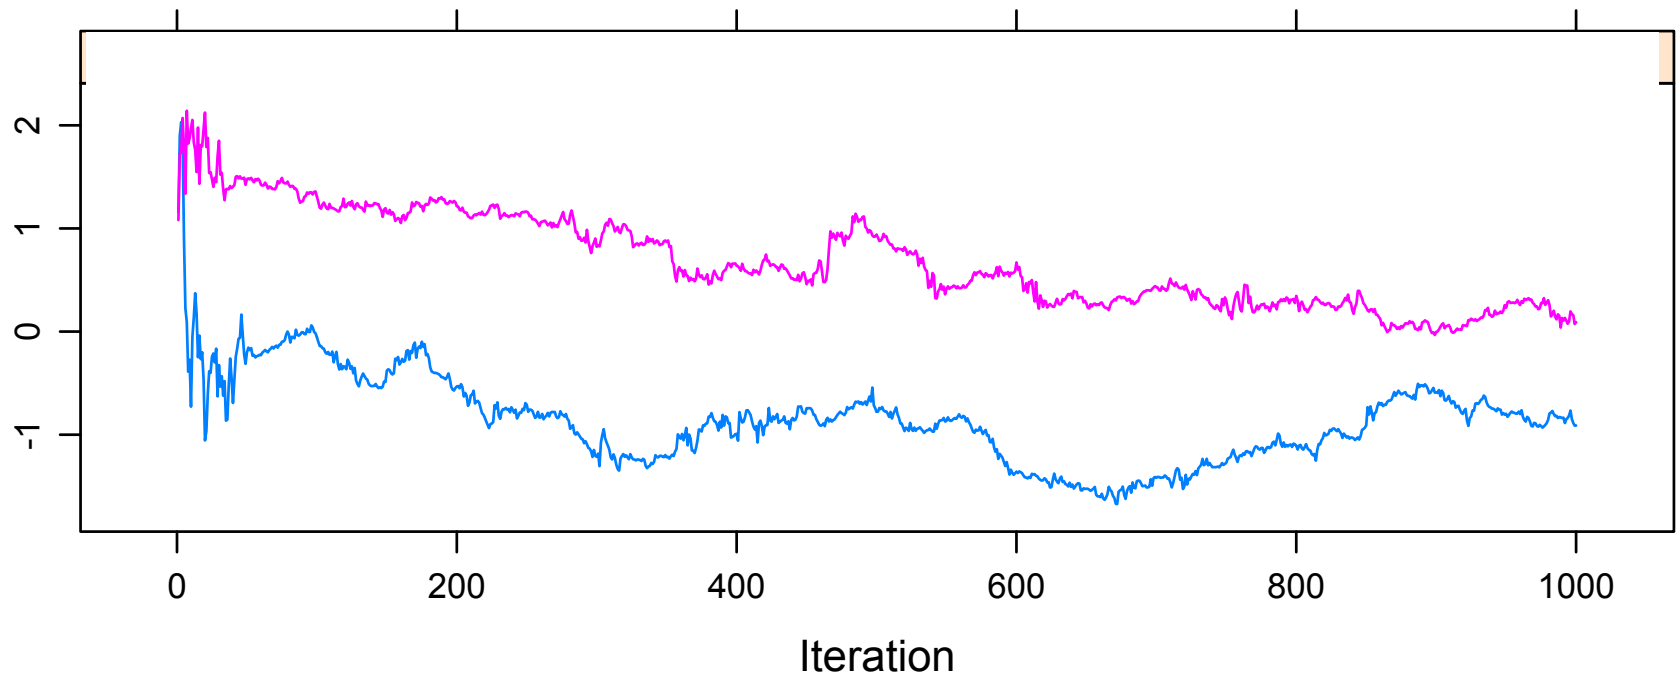
Have these two chains (pink and blue) converged?



**Looks good**  
**[We are OK to look at the model results]**

# Visual assessment of convergence

Have these two chains (pink and blue) converged?



**Looks bad**  
**[Model results cannot be used]**

# A basic example: prevalence estimation

# Define and run the model

Define the model, load the runjags library, and define the data in R:

```
bugs_model <- "
  model{
    Pos ~ dbinom(ap, N)
    # Uniform (non-informative) prior for apparent prev.
    ap ~ dbeta(1,1)

    #data# Pos, N
    #inits# ap
    #monitor# ap
  }
"

library('runjags')

# Data:
Pos <- 1210
N <- 4072

# Initial values:
ap <- list(chain1=0.1, chain2=0.9)

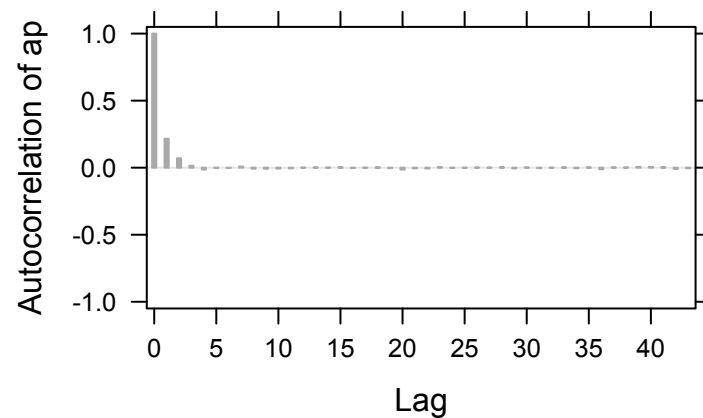
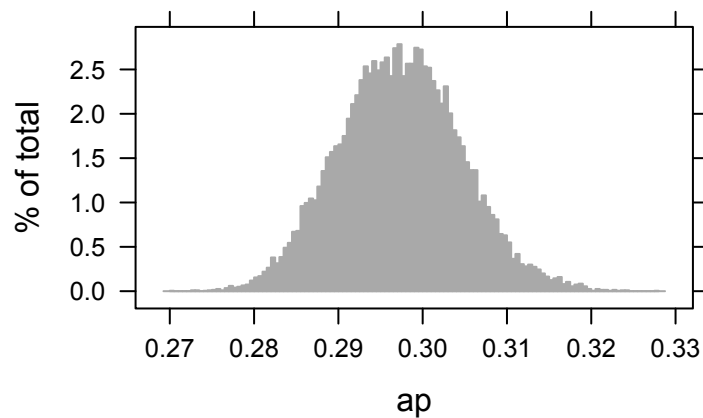
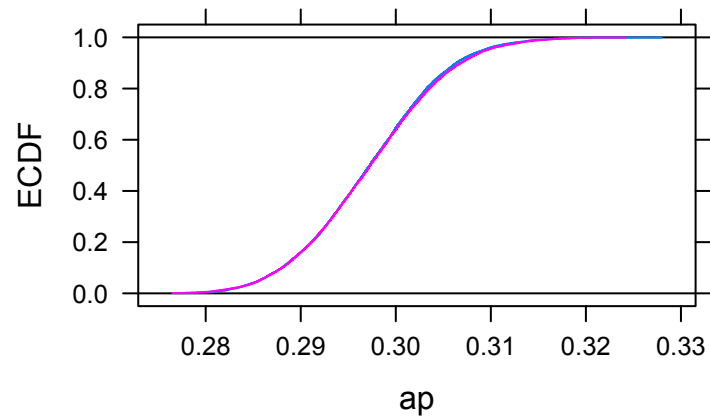
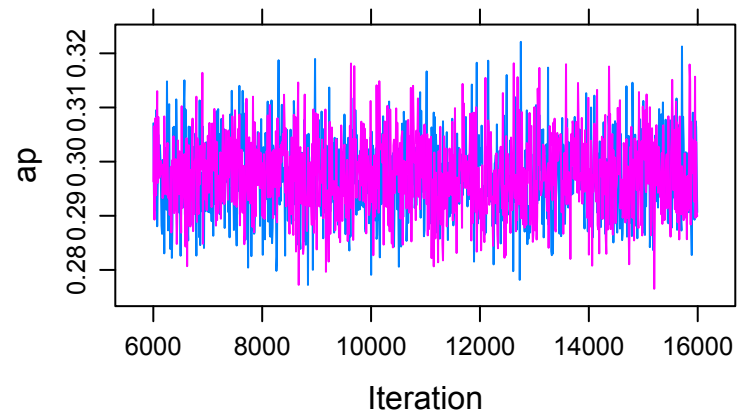
results <- run.jags(bugs_model, n.chains=2, burnin=5000, sample=10000)
```





# Check convergence

`plot(results)`



# Check sample size and results

```
> results
```

JAGS model summary statistics from 20000 samples (chains = 2;  
adapt+burnin = 6000):

|    | Lower95 | Median  | Upper95 | Mean   | SD        | Mode |
|----|---------|---------|---------|--------|-----------|------|
| ap | 0.28339 | 0.29719 | 0.31155 | 0.2972 | 0.0071436 | --   |

|    | MCerr       | MC%ofSD | SSeff | AC.10    | psrf   |
|----|-------------|---------|-------|----------|--------|
| ap | 0.000063509 | 0.9     | 12652 | 0.018843 | 1.0002 |

Total time taken: 0.8 seconds

# See also:

?runjags

